

Artificial Intelligence (CS303)

Lecture 3: Problem-Specific Search

Hints for this lecture

- The more we know about the problem characteristic/structure, the better we can solve it.

Outline of this lecture

- **Make Search Algorithms Less General**
- **Gradient-based Methods for Numerical Optimization**
- **Quadratic Programming Problems**
- **Constraint Satisfaction Problems**
- **Adversarial Search**


Make Search Algorithms Less General

- The search methods talked previously are rather general, i.e., applicable to any problem.
- Generality is nice and worthy of pursuit, while it usually conflicts with the other desired features of algorithms, e.g., efficiency.
- A search method is general because the characteristic/structure (no matter we know or not) of the problem **is not taken into account** when designing the search method.

Make Search Algorithms Less General

- When designing an algorithm for a problem (class), taking the problem characteristics into account usually helps us get the desired solution by **searching only a part of the search/state space**, making the search more efficient.

- In some cases, we do know something about problem.



We know all previous steps

- As long as a problem (class) is **of sufficient significance**, it is worthy of designing problem-specific algorithm for it.

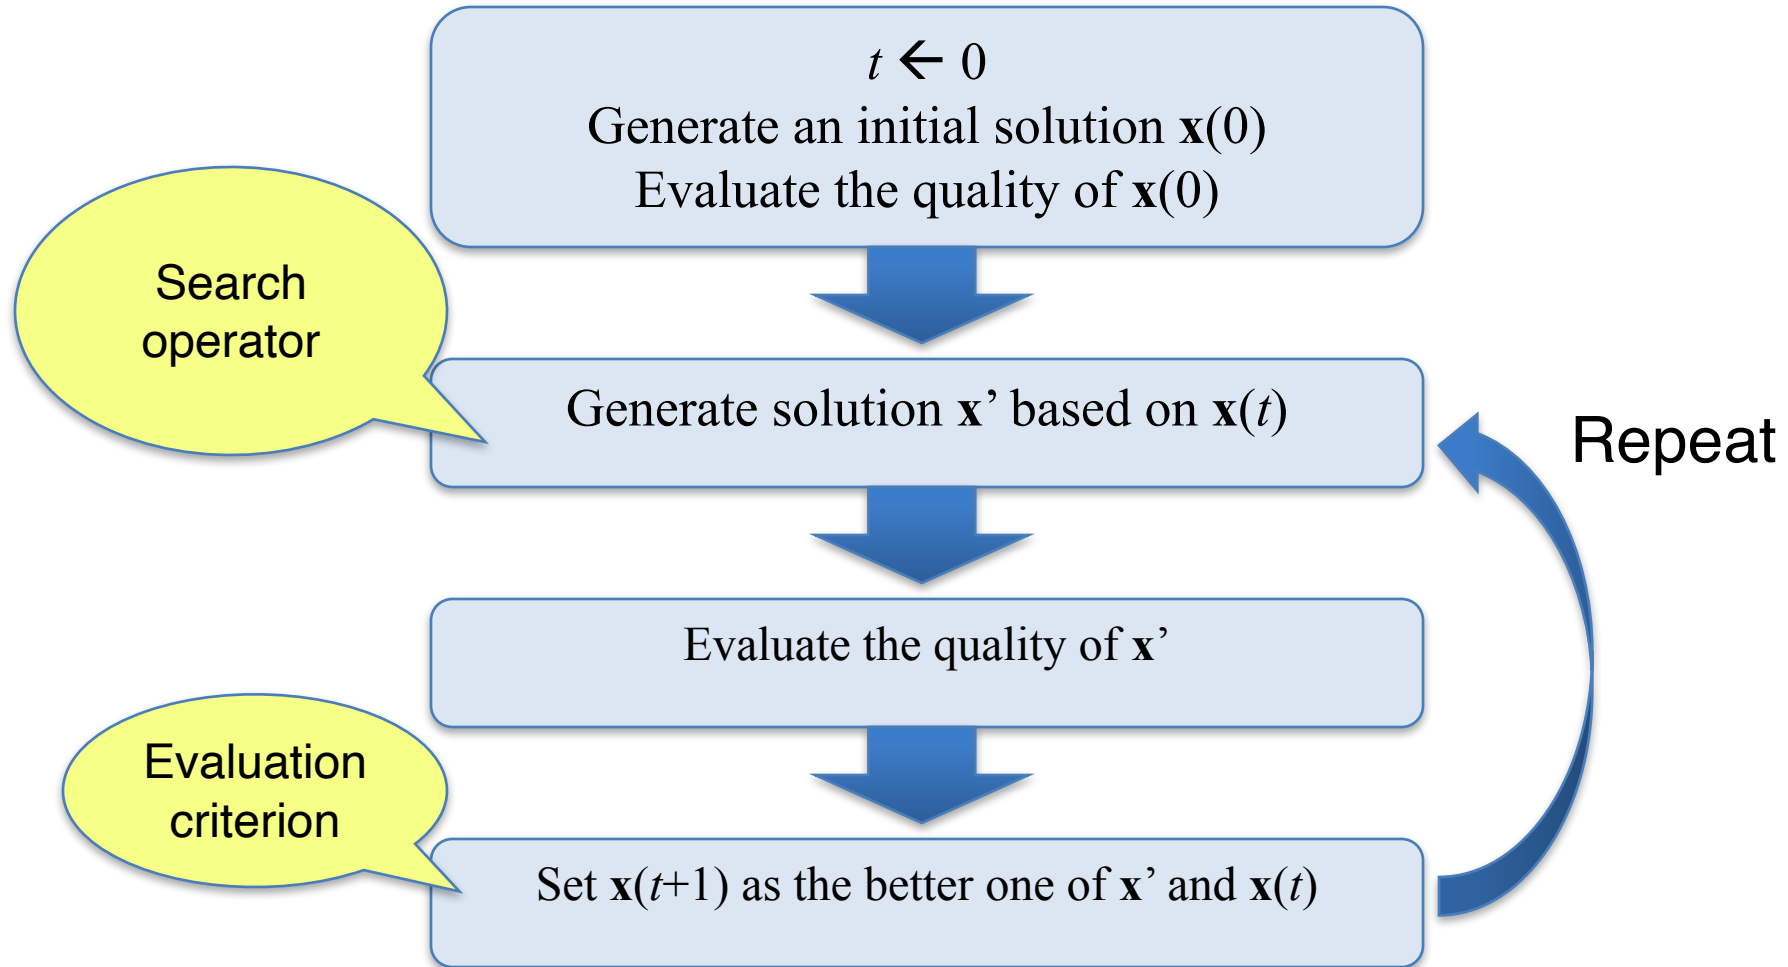
Make Search Algorithms Less General

- Again, consider the ubiquitous optimization problems.

$$\begin{aligned} &\text{maximize } f(x) \\ &\text{subject to: } g_i(x) \leq 0, \quad i = 1 \dots m \\ &\quad \quad \quad h_j(x) = 0, \quad j = 1 \dots p \end{aligned}$$

- What do you mean by “problem characteristic”? Most basically:
 - What is x ?
 - What is f ?
 - Does f fulfill some properties that would lead to a more efficient search?

Recall The General Framework for Search



Gradient-based Methods for Numerical Optimization

- Suppose the objective function $f(x_1, y_1, x_2, y_2, x_3, y_3)$ is **continuous and differentiable** (thus the gradient could be calculated)

Gradient methods compute


$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial y_1}, \frac{\partial f}{\partial x_2}, \frac{\partial f}{\partial y_2}, \frac{\partial f}{\partial x_3}, \frac{\partial f}{\partial y_3} \right)$$

to increase/reduce f , e.g., by $\mathbf{x} \leftarrow \mathbf{x} + \alpha \nabla f(\mathbf{x})$

Quadratic Programming Problems

- The objective function is a **quadratic** function of x
 - stronger condition than just differentiable.
- The constraints are linear functions of x

$$\begin{aligned} &\text{maximize } f(x) \\ &\text{subject to: } g_i(x) \leq 0, \quad i = 1 \dots m \\ &\quad \quad \quad h_j(x) = 0, \quad j = 1 \dots p \end{aligned}$$



$$\begin{aligned} &\min f(x) = q^T x + \frac{1}{2} x^T Q x \\ &s.t. Ax = a \\ &\quad Bx \leq b \\ &\quad x \geq 0 \end{aligned}$$

Quadratic Programming Problems

- We take an even stronger condition as example
 - no constraints.
 - The objective function is not only quadratic, but also **convex**.
 - $f(x)$ is a convex function of x

$$\min f(x) = q^T x + \frac{1}{2} x^T Q x$$

$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y).$$





Figure 1: An illustration of the definition of a convex function

Quadratic Programming Problems

- How to solve such a problem by search?
 - Simply set the derivative of f to 0, and solve a linear system
 - No need to *search* at all!

$$\min f(x) = q^T x + \frac{1}{2} x^T Q x$$



- More practical cases still needs search (e.g., conjugate gradient method for QP with (e.g., with constraints), recall the Lagrange multiplier technique.

Quadratic Programming Problems

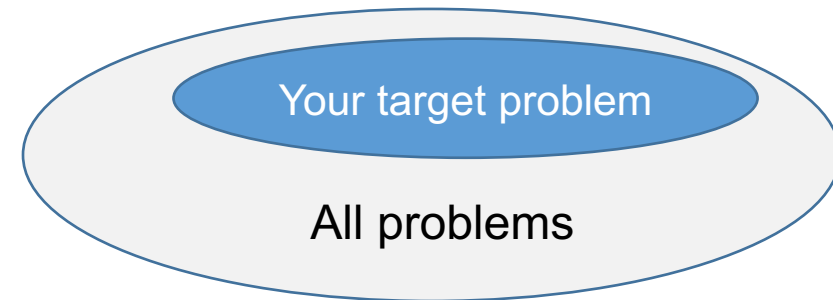
- How do I know the objective function is convex?
 - A sufficient condition: Q is positive definite.

$$\begin{aligned} & q^T [\lambda x + (1-\lambda)y] + \frac{1}{2} [\lambda x + (1-\lambda)y]^T Q [\lambda x + (1-\lambda)y] \\ &= \lambda q^T x + (1-\lambda) q^T y + \frac{1}{2} \lambda^2 x^T Q x + \frac{1}{2} (1-\lambda)^2 y^T Q y \\ &\quad + \frac{1}{2} \lambda (1-\lambda) x^T Q y + \frac{1}{2} \lambda (1-\lambda) y^T Q x \\ &\leq \\ & \lambda q^T x + \frac{1}{2} \lambda x^T Q x + (1-\lambda) q^T y + \frac{1}{2} (1-\lambda) y^T Q y \end{aligned}$$

Quadratic Programming Problems

Lesson learned from the simple example

- if the problem have very good property, we can even reduce the search process to a single step (solve analytically).
- Needs to carefully check whether the “good property” holds.
- Intuitively, better property corresponds to stronger conditions
 - more unlikely to hold
 - application-domain of search algorithm developed based on such properties is more restrictive.



Constraint Satisfaction Problems

Standard search problem:

state is a “black box”—any old data structure
that supports goal test, eval, successor

CSP:

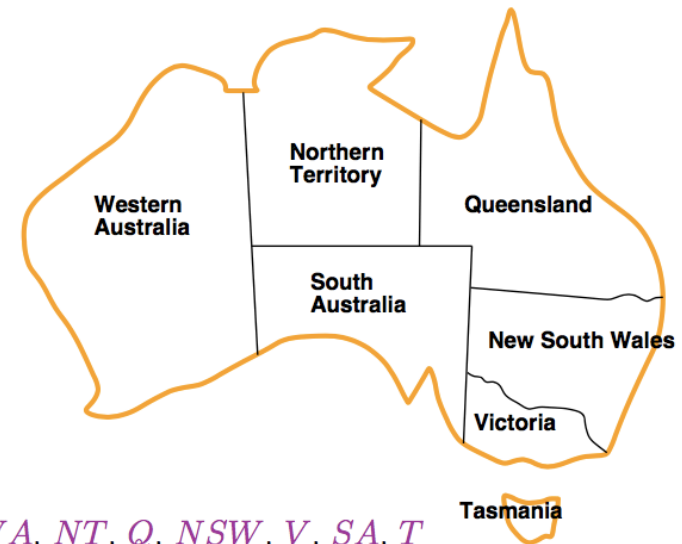
state is defined by **variables** X_i with **values** from domain D_i

goal test is a set of **constraints** specifying
allowable combinations of values for subsets of variables

Simple example of a **formal representation language**

Allows useful **general-purpose** algorithms with more power
than standard search algorithms

Example: Map Coloring



Variables WA, NT, Q, NSW, V, SA, T

Domains $D_i = \{red, green, blue\}$

Constraints: adjacent regions must have different colors

e.g., $WA \neq NT$ (if the language allows this), or

$(WA, NT) \in \{(red, green), (red, blue), (green, red), (green, blue), \dots\}$

Variants of CSPs

- Unary constraints involve a single variable.
- Binary constraints involve pairs of variables.
- Higher-order constraints involve 3 or more variables.
- Preferences (Soft constraints), e.g., **red** is better than **green**, is often represented by a cost for each variable assignment (i.e., the target is to minimize the cost).

Real-world CSPs

- Assignment problems
- Timetabling problems
- Hardware configuration
- Floorplanning
- Factory scheduling
- ...

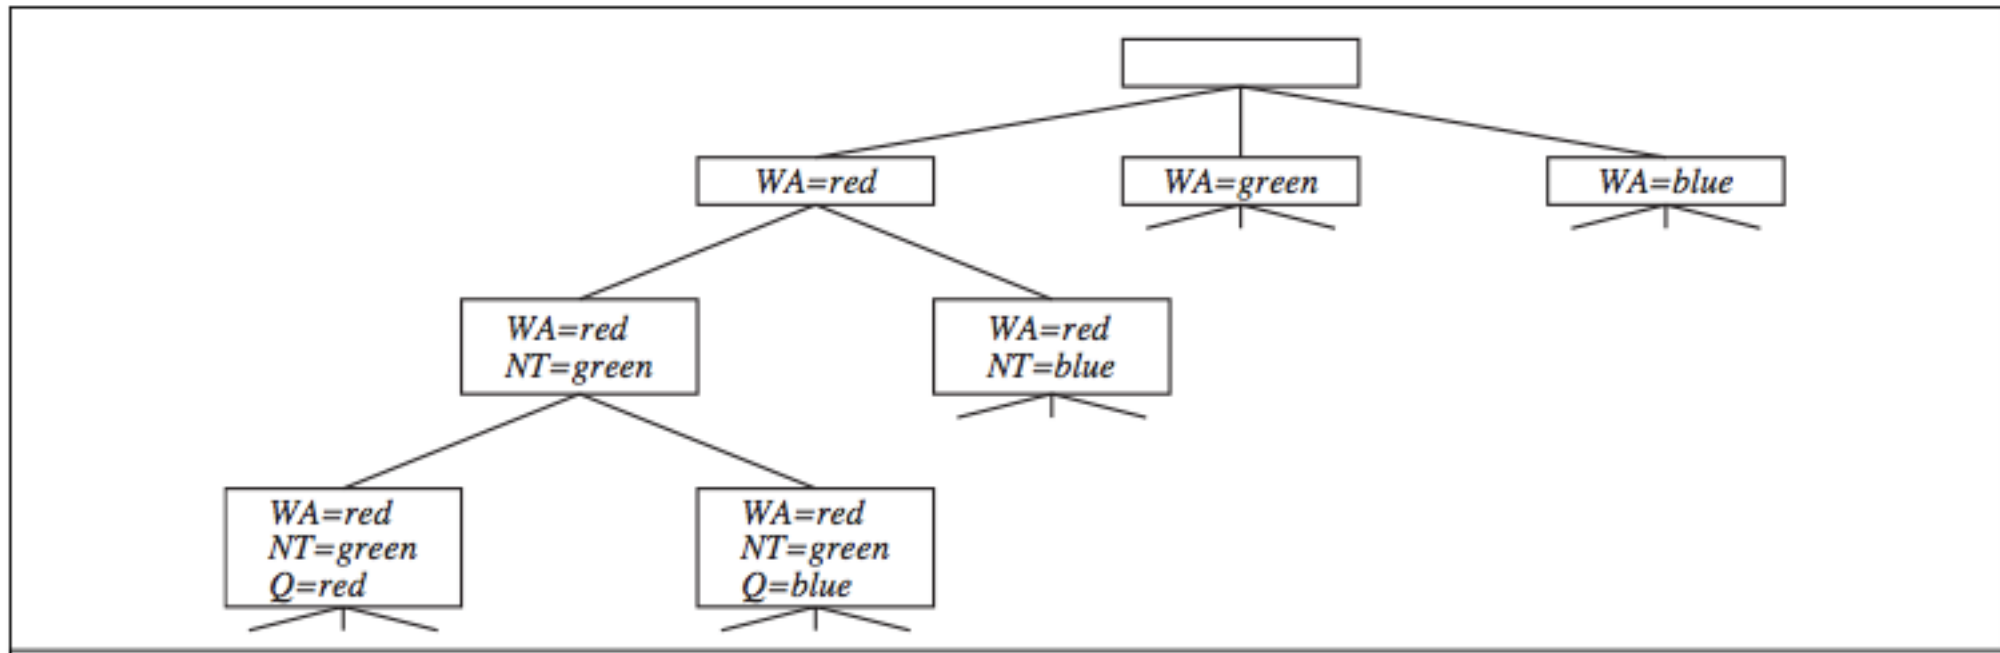
What is the search tree of a CSP?

Characteristics of CSPs

- **Commutativity:** the order of assigning values to variables does not affect the final outcome.
- **The constraints provide additional information that could be represented by a constraint graph.**

Commutativity


- Commutativity help us formulate the search tree (only 1 variable needs to be considered at each node in the search tree).



Constraint Graph

Binary CSP: each constraint relates at most two variables

Constraint graph: nodes are variables, arcs show constraints



General-purpose CSP algorithms use the graph structure to speed up search. E.g., Tasmania is an independent subproblem!

Inference

- A constraint graph allows the agent to do **inference** in addition to search.
- Inference basically means **checking local consistency** (or detecting inconsistency)
 - Node consistency
 - Arc Consistency
 - Path Consistency
 - K -consistency
 - Global consistency
- Inference helps **prune** the search tree, either before or during the search.

Backtracking Search for CSP

- Depth-first search, assign a value to unassigned variables recursively.
- If inconsistency occurred, move 1 step back to try another value.

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

Improving Backtracking Search (1)

Applying inference

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

Applying inference

Improving Backtracking Search (2)

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

Choosing variables
with minimum
numbers of
remaining value

Improving Backtracking Search (3)


```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

Maintain a conflict
set and do
backjumping



Local Search for CSP

- CSP can be actually reformulated as a **constraint optimization problem**, for which the objective function is to minimize the constraint violation.
- Working in the solution space (complete solution formulation)
- Iteratively select a conflicted variable and assign a different value to it.
- Choose the value that leads to the minimum cost.



Games as A Search Problem


- More than 1 agent/player, with conflicting goals.
- We consider two players, zero-sum games.



Definition of a game

- S_0 : The **initial state**, which specifies how the game is set up at the start.
- $\text{PLAYER}(s)$: Defines which player has the move in a state.
- $\text{ACTIONS}(s)$: Returns the set of legal moves in a state.
- $\text{RESULT}(s, a)$: The **transition model**, which defines the result of a move.
- $\text{TERMINAL-TEST}(s)$: A **terminal test**, which is true when the game is over and false otherwise. States where the game has ended are called **terminal states**.
- $\text{UTILITY}(s, p)$: A **utility function** (also called an objective function or payoff function),


Tic-Tac-Toe Search Tree



Two players: MAX and MIN

Optimal decision in games

- Assume the game is deterministic and perfect information is available
- Idea (for MAX): choose the move to position **the highest minimax value**



MINIMAX(s) =


$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Minimax Algorithm

- Perform a complete depth-first search of the game tree.
- Recursively compute the minimax values of each successor state.
- Maximize the worst-case outcome for MAX.

Alpha-Beta Pruning

- A simplification of minimax algorithm.
- Remove (unneeded) part of the minimax tree from consideration.



Alpha: the value of the best (i.e., **highest**-value) choice we have found so far at any choice point along the path for **MAX**.

Beta: the value of the best (i.e., **lowest**-value) choice we have found so far at any choice point along the path for **MIN**.

Alpha-Beta Pruning



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return v  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return v  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return v
```

The search order is important

- It might be worthwhile to examine first the successors that are likely to be the best.



Further acceleration?

- Alpha-beta pruning does **not** affect the final result. (good!)
- With perfect ordering, the time complexity is still high, i.e., $O(b^{m/2})$. (bad...)

In practice:

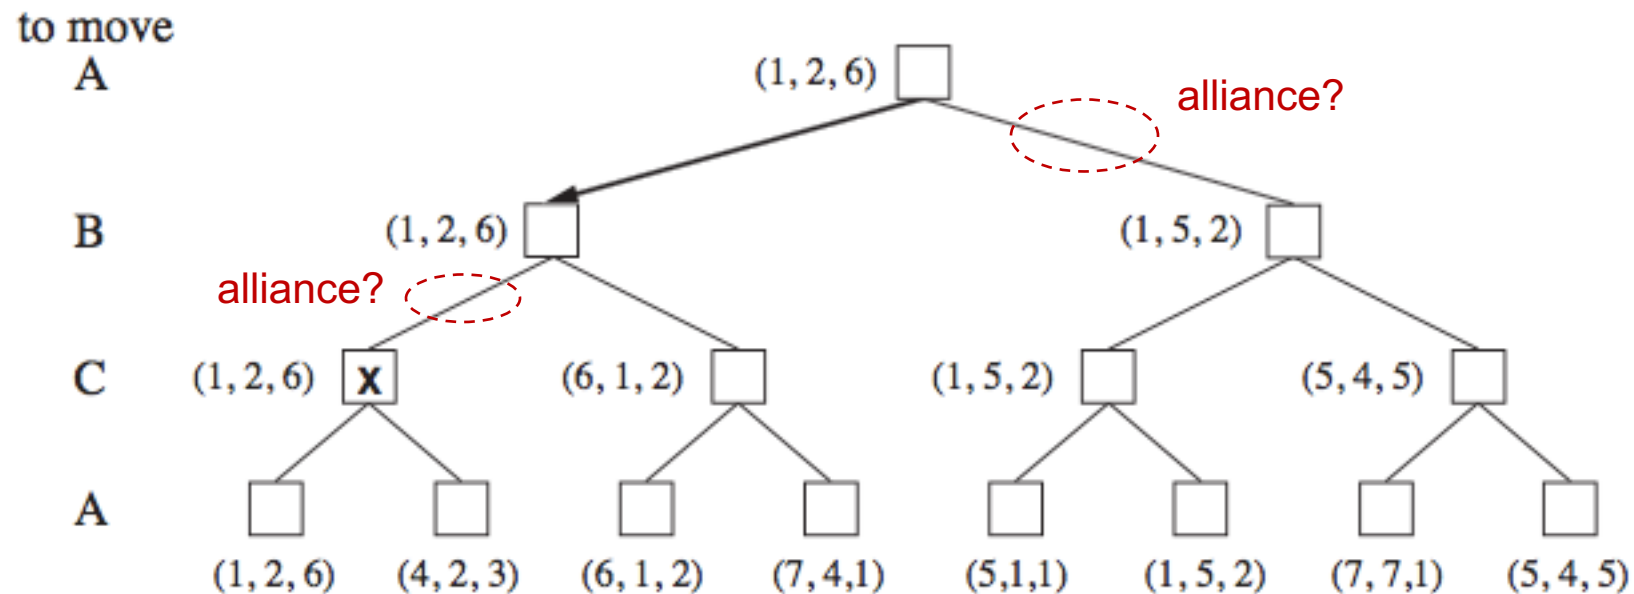
- Replace the **UTILITY** function with a heuristic evaluation function **EVAL**
- Use **CUTOFF-Test** instead of **TERMINAL-Test**

H-MINIMAX(s, d) =

$$\begin{cases} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MIN.} \end{cases}$$

More Complex Situations

- games with more than 2 players
- 2-players game that is not zero-sum
- Minimax or Alpha-Beta Pruning don't apply



Question to ask when tackling a search problem

- How to **represent** the search space?
 - Search Tree (state space)
 - Solution space
- What is the **objective function and constraint**, and algorithm in textbook already good enough?
- Which **algorithmic framework** to choose?
 - Tree search, e.g., Un-informed Search, Heuristic Search (A*...)
 - Direct search in the solution space, e.g., Hill Climbing, Simulated Annealing, Genetic Algorithm...
- How to define **concrete components** of the algorithm framework?
 - General-purpose operators in literature
 - Problem-specific operators, designed based on domain knowledge

Always **trade-off** among solution quality, efficiency, and your domain knowledge

The End of The **Search** Section.